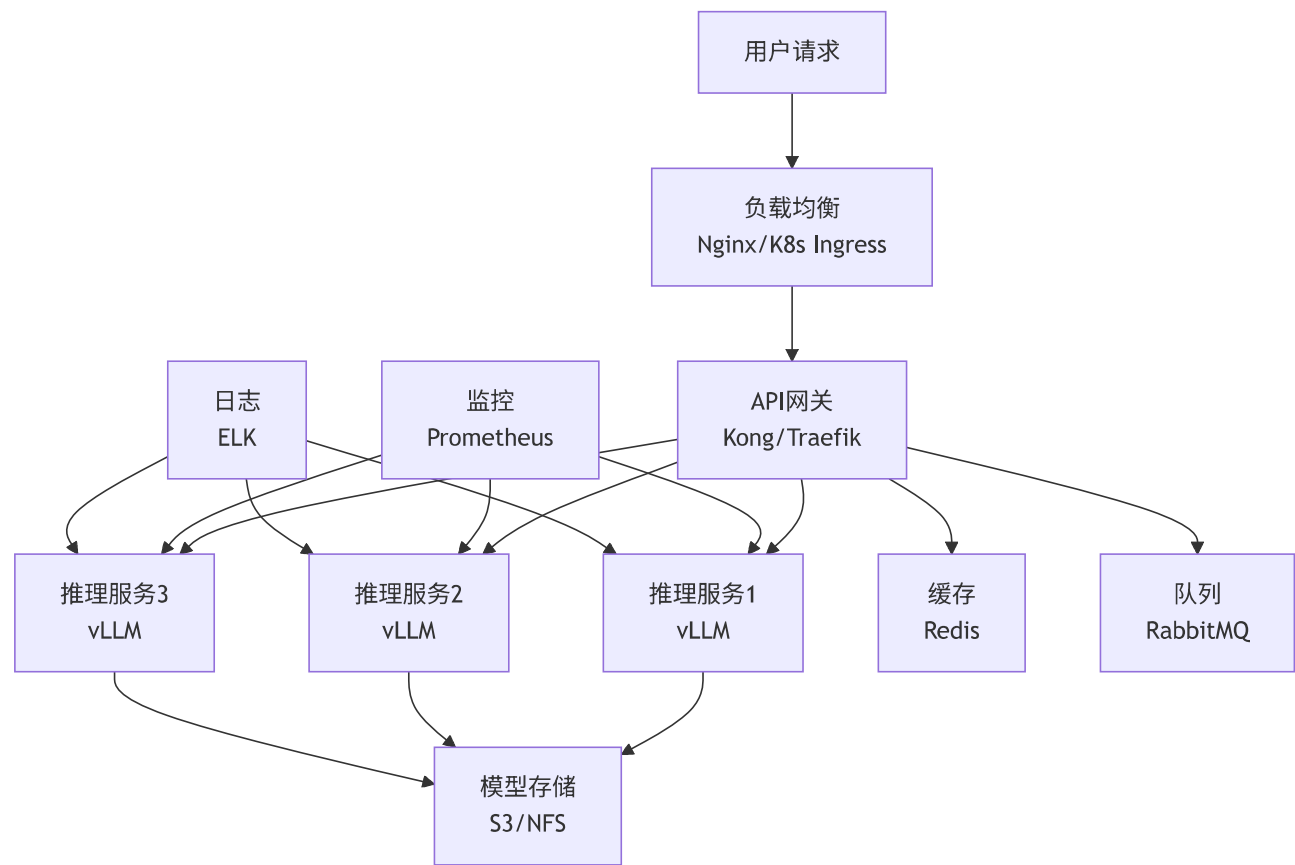


第10章 部署与服务化

将大模型从实验室搬到生产环境

10.1 部署架构概览

10.1.1 典型部署架构



10.1.2 部署方式对比

方式	优点	缺点	适用场景
本地部署	完全控制、低延迟	需要GPU资源	开发测试
云服务API	简单、按需付费	成本高、依赖外部	快速原型
自建GPU集群	灵活、成本可控	运维复杂	生产环境
Serverless	自动扩展、按用量	冷启动慢	间歇性负载
边缘部署	低延迟、隐私	资源受限	IoT设备

10.2 使用vLLM部署

10.2.1 基础部署

```

# 1. 安装
pip install vllm

# 2. Python API
from vllm import LLM, SamplingParams

# 初始化模型
llm = LLM(
    model="meta-llama/Llama-2-7b-chat-hf",
    tensor_parallel_size=2, # 2张GPU并行
    dtype="float16",
    max_model_len=4096
)

# 生成
prompts = [
    "The capital of France is",
    "The largest planet is"
]

sampling_params = SamplingParams(
    temperature=0.8,
    top_p=0.95,
    max_tokens=256
)

outputs = llm.generate(prompts, sampling_params)

for output in outputs:
    print(output.outputs[0].text)

```

10.2.2 OpenAI兼容服务器

```

# 启动服务器
python -m vllm.entrypoints.openai.api_server \
    --model meta-llama/Llama-2-7b-chat-hf \
    --tensor-parallel-size 2 \
    --port 8000

# 客户端调用（兼容OpenAI API）
from openai import OpenAI

client = OpenAI(
    base_url="http://localhost:8000/v1",
    api_key="fake-key" # vLLM不需要真实key

```

```
)

response = client.chat.completions.create(
    model="meta-llama/Llama-2-7b-chat-hf",
    messages=[
        {"role": "user", "content": "你好，介绍一下自己"}
    ]
)

print(response.choices[0].message.content)
```

10.2.3 Docker部署

```
# Dockerfile
FROM vllm/vllm-openai:latest

# 设置环境变量
ENV MODEL_NAME=meta-llama/Llama-2-7b-chat-hf
ENV TENSOR_PARALLEL_SIZE=2
ENV GPU_MEMORY_UTILIZATION=0.9

# 启动命令
CMD python -m vllm.entrypoints.openai.api_server \
    --model $MODEL_NAME \
    --tensor-parallel-size $TENSOR_PARALLEL_SIZE \
    --gpu-memory-utilization $GPU_MEMORY_UTILIZATION \
    --host 0.0.0.0 \
    --port 8000

# 构建镜像
docker build -t my-llm-service .

# 运行容器
docker run --gpus all \
    -p 8000:8000 \
    -v ~/.cache/huggingface:/root/.cache/huggingface \
    my-llm-service
```

10.3 API服务设计

10.3.1 FastAPI实现

```

from fastapi import FastAPI, HTTPException
from pydantic import BaseModel
from typing import List, Optional
import asyncio

app = FastAPI(title="LLM API Service")

# 请求模型
class CompletionRequest(BaseModel):
    prompt: str
    max_tokens: int = 256
    temperature: float = 0.8
    top_p: float = 0.95
    n: int = 1
    stream: bool = False

class ChatMessage(BaseModel):
    role: str
    content: str

class ChatRequest(BaseModel):
    messages: List[ChatMessage]
    max_tokens: int = 256
    temperature: float = 0.8
    top_p: float = 0.95
    stream: bool = False

# 初始化模型（全局单例）
from vllm import LLM, SamplingParams

llm = LLM(model="meta-llama/Llama-2-7b-chat-hf")

@app.post("/v1/completions")
async def create_completion(request: CompletionRequest):
    """
    文本补全接口
    """
    try:
        sampling_params = SamplingParams(
            temperature=request.temperature,
            top_p=request.top_p,
            max_tokens=request.max_tokens,
            n=request.n
        )

        outputs = llm.generate([request.prompt], sampling_params)

```

```

        return {
            "choices": [
                {
                    "text": output.outputs[0].text,
                    "index": 0,
                    "finish_reason": "stop"
                }
            ]
        }
    except Exception as e:
        raise HTTPException(status_code=500, detail=str(e))

```

```

@app.post("/v1/chat/completions")
async def create_chat_completion(request: ChatRequest):
    """
    对话接口
    """
    try:
        # 构造prompt (根据模型格式)
        prompt = format_chat_prompt(request.messages)

        sampling_params = SamplingParams(
            temperature=request.temperature,
            top_p=request.top_p,
            max_tokens=request.max_tokens
        )

        if request.stream:
            # 流式输出
            return StreamingResponse(
                stream_generator(prompt, sampling_params),
                media_type="text/event-stream"
            )
        else:
            # 普通输出
            outputs = llm.generate([prompt], sampling_params)

            return {
                "choices": [{
                    "message": {
                        "role": "assistant",
                        "content": outputs[0].outputs[0].text
                    },

```

```

        "finish_reason": "stop"
    }]
}
except Exception as e:
    raise HTTPException(status_code=500, detail=str(e))

def format_chat_prompt(messages: List[ChatMessage]) -> str:
    """
    格式化对话为prompt (LLaMA-2格式)
    """
    prompt = "<s>"
    for msg in messages:
        if msg.role == "system":
            prompt += f"[INST] <<SYS>>\n{msg.content}\n</SYS>>\n\n"
        elif msg.role == "user":
            prompt += f"[INST] {msg.content} [/INST] "
        elif msg.role == "assistant":
            prompt += f"{msg.content} </s><s>"

    return prompt

async def stream_generator(prompt, sampling_params):
    """
    流式生成器
    """
    # 使用异步生成
    async for text in llm.generate_async(prompt, sampling_params):
        yield f"data: {json.dumps({'text': text})}\n\n"

    yield "data: [DONE]\n\n"

@app.get("/health")
async def health_check():
    """
    健康检查接口
    """
    return {"status": "ok"}

# 启动服务
if __name__ == "__main__":
    import uvicorn
    uvicorn.run(app, host="0.0.0.0", port=8000, workers=1)

```

10.3.2 限流和缓存

```
from fastapi import FastAPI, Request
from slowapi import Limiter, _rate_limit_exceeded_handler
from slowapi.util import get_remote_address
from slowapi.errors import RateLimitExceeded
import redis
import hashlib

app = FastAPI()

# 限流器
limiter = Limiter(key_func=get_remote_address)
app.state.limiter = limiter
app.add_exception_handler(RateLimitExceeded, _rate_limit_exceeded_handler)

# Redis缓存
redis_client = redis.Redis(host='localhost', port=6379, db=0)

def get_cache_key(prompt: str, params: dict) -> str:
    """
    生成缓存key
    """
    content = f"{prompt}_{json.dumps(params, sort_keys=True)}"
    return hashlib.md5(content.encode()).hexdigest()

@app.post("/v1/completions")
@limiter.limit("100/minute") # 每分钟最多100次请求
async def create_completion(request: Request, req: CompletionRequest):
    """
    带缓存的补全接口
    """
    # 1. 检查缓存
    cache_key = get_cache_key(req.prompt, {
        "temperature": req.temperature,
        "max_tokens": req.max_tokens
    })

    cached = redis_client.get(cache_key)
    if cached:
        return json.loads(cached)

    # 2. 生成结果
    result = await generate_completion(req)
```

```
# 3. 存入缓存（过期时间1小时）
redis_client.setex(cache_key, 3600, json.dumps(result))

return result
```

10.4 批处理和队列

10.4.1 异步批处理

```
import asyncio
from collections import deque

class BatchProcessor:
    """
    批处理器：收集请求，批量推理
    """

    def __init__(self, model, max_batch_size=32, max_wait_time=0.1):
        self.model = model
        self.max_batch_size = max_batch_size
        self.max_wait_time = max_wait_time

        self.queue = deque()
        self.processing = False

    async def add_request(self, prompt, params):
        """
        添加请求到队列
        """
        future = asyncio.Future()
        self.queue.append((prompt, params, future))

        # 触发处理
        if not self.processing:
            asyncio.create_task(self.process_batch())

        # 等待结果
        return await future

    async def process_batch(self):
        """
        处理一批请求
        """
        self.processing = True

        # 等待accumulate
```



```

await asyncio.sleep(self.max_wait_time)

# 收集batch
batch = []
futures = []

while self.queue and len(batch) < self.max_batch_size:
    prompt, params, future = self.queue.popleft()
    batch.append((prompt, params))
    futures.append(future)

if batch:
    # 批量推理
    prompts = [item[0] for item in batch]
    results = self.model.generate(prompts)

    # 返回结果给各个请求
    for future, result in zip(futures, results):
        future.set_result(result)

self.processing = False

# 如果还有请求，继续处理
if self.queue:
    asyncio.create_task(self.process_batch())

# 在FastAPI中使用
batch_processor = BatchProcessor(llm)

@app.post("/v1/completions")
async def create_completion(request: CompletionRequest):
    result = await batch_processor.add_request(
        request.prompt,
        {"temperature": request.temperature}
    )
    return result

```

10.4.2 任务队列 (Celery)

```

from celery import Celery
import redis

# 初始化Celery
celery_app = Celery(

```

```

        'llm_tasks',
        broker='redis://localhost:6379/0',
        backend='redis://localhost:6379/1'
    )

@celery_app.task(bind=True)
def generate_text_task(self, prompt, params):
    """
    异步生成任务
    """
    try:
        # 更新任务状态
        self.update_state(state='PROCESSING')

        # 生成
        result = llm.generate([prompt], SamplingParams(**params))

        return {
            "status": "success",
            "text": result[0].outputs[0].text
        }
    except Exception as e:
        return {
            "status": "error",
            "error": str(e)
        }

# FastAPI接口
@app.post("/v1/async/completions")
async def create_async_completion(request: CompletionRequest):
    """
    异步任务接口
    """
    task = generate_text_task.delay(
        request.prompt,
        {
            "temperature": request.temperature,
            "max_tokens": request.max_tokens
        }
    )

    return {
        "task_id": task.id,
        "status": "queued"
    }

```

```

@app.get("/v1/async/completions/{task_id}")
async def get_task_result(task_id: str):
    """
    查询任务结果
    """

    task = generate_text_task.AsyncResult(task_id)

    if task.ready():
        return {
            "status": "completed",
            "result": task.result
        }
    else:
        return {
            "status": "processing"
        }

```

10.5 监控和日志

10.5.1 Prometheus监控

```

from prometheus_client import Counter, Histogram, Gauge, generate_latest
from fastapi import Response
import time

# 定义指标
request_count = Counter(
    'llm_requests_total',
    'Total number of requests',
    ['endpoint', 'status']
)

request_duration = Histogram(
    'llm_request_duration_seconds',
    'Request duration in seconds',
    ['endpoint']
)

model_gpu_memory = Gauge(
    'llm_gpu_memory_used_bytes',
    'GPU memory used in bytes',
    ['gpu_id']
)

```

```
active_requests = Gauge(  
    'llm_active_requests',  
    'Number of active requests'  
)
```

```
# 中间件
```

```
@app.middleware("http")
```

```
async def monitor_middleware(request: Request, call_next):
```

```
    """
```

```
    监控中间件
```

```
    """
```

```
    active_requests.inc()
```

```
    start_time = time.time()
```

```
    try:
```

```
        response = await call_next(request)
```

```
        # 记录指标
```

```
        duration = time.time() - start_time
```

```
        request_duration.labels(endpoint=request.url.path).observe(duration)
```

```
        request_count.labels(  
            endpoint=request.url.path,
```

```
            status=response.status_code  
        ).inc()
```

```
        return response
```

```
    finally:
```

```
        active_requests.dec()
```

```
# 暴露metrics
```

```
@app.get("/metrics")
```

```
async def metrics():
```

```
    """
```

```
    Prometheus metrics endpoint
```

```
    """
```

```
    return Response(  
        content=generate_latest(),  
        media_type="text/plain"
```

```
    )
```

```
# 更新GPU内存指标（定时任务）
```

```
import torch
```

```

def update_gpu_metrics():
    """
    更新GPU metrics
    """
    for i in range(torch.cuda.device_count()):
        memory = torch.cuda.memory_allocated(i)
        model_gpu_memory.labels(gpu_id=str(i)).set(memory)

```

10.5.2 结构化日志

```

import logging
import json
from pythonjsonlogger import jsonlogger

# 配置JSON日志
logHandler = logging.StreamHandler()
formatter = jsonlogger.JsonFormatter(
    '%(asctime)s %(name)s %(levelname)s %(message)s'
)
logHandler.setFormatter(formatter)

logger = logging.getLogger()
logger.addHandler(logHandler)
logger.setLevel(logging.INFO)

# 在API中使用
@app.post("/v1/completions")
async def create_completion(request: CompletionRequest):
    request_id = str(uuid.uuid4())

    logger.info(
        "Request received",
        extra={
            "request_id": request_id,
            "prompt_length": len(request.prompt),
            "temperature": request.temperature,
            "max_tokens": request.max_tokens
        }
    )

    start_time = time.time()

    try:

```

```

        result = await generate_completion(request)

        logger.info(
            "Request completed",
            extra={
                "request_id": request_id,
                "duration": time.time() - start_time,
                "tokens_generated": len(result["choices"][0]["text"].split())
            }
        )

        return result
    except Exception as e:
        logger.error(
            "Request failed",
            extra={
                "request_id": request_id,
                "error": str(e),
                "duration": time.time() - start_time
            },
            exc_info=True
        )
        raise

```

10.6 Kubernetes部署

10.6.1 Deployment配置

```

# llm-deployment.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: llm-service
  labels:
    app: llm-service
spec:
  replicas: 3
  selector:
    matchLabels:
      app: llm-service
  template:
    metadata:
      labels:
        app: llm-service
    spec:

```

```

containers:
- name: llm-service
  image: my-llm-service:latest
  ports:
  - containerPort: 8000
  resources:
    requests:
      nvidia.com/gpu: 1
      memory: "16Gi"
      cpu: "4"
    limits:
      nvidia.com/gpu: 1
      memory: "32Gi"
      cpu: "8"
  env:
  - name: MODEL_NAME
    value: "meta-llama/Llama-2-7b-chat-hf"
  - name: TENSOR_PARALLEL_SIZE
    value: "1"
  volumeMounts:
  - name: model-cache
    mountPath: /root/.cache/huggingface
volumes:
- name: model-cache
  persistentVolumeClaim:
    claimName: model-cache-pvc

```

10.6.2 Service和Ingress

```

# Llm-service.yaml
apiVersion: v1
kind: Service
metadata:
  name: llm-service
spec:
  selector:
    app: llm-service
  ports:
  - port: 8000
    targetPort: 8000
  type: ClusterIP

```

```

# Llm-ingress.yaml
apiVersion: networking.k8s.io/v1

```

```

kind: Ingress
metadata:
  name: llm-ingress
  annotations:
    nginx.ingress.kubernetes.io/proxy-body-size: "10m"
    nginx.ingress.kubernetes.io/proxy-read-timeout: "300"
spec:
  rules:
  - host: llm.example.com
    http:
      paths:
      - path: /
        pathType: Prefix
        backend:
          service:
            name: llm-service
            port:
              number: 8000

```

10.6.3 HPA (水平自动扩缩容)

```

# Llm-hpa.yaml
apiVersion: autoscaling/v2
kind: HorizontalPodAutoscaler
metadata:
  name: llm-hpa
spec:
  scaleTargetRef:
    apiVersion: apps/v1
    kind: Deployment
    name: llm-service
  minReplicas: 2
  maxReplicas: 10
  metrics:
  - type: Resource
    resource:
      name: cpu
      target:
        type: Utilization
        averageUtilization: 70
  - type: Pods
    pods:
      metric:
        name: llm_active_requests
      target:

```



```
type: AverageValue
averageValue: "50"
```

10.7 性能优化

10.7.1 性能指标

指标	定义	目标
延迟 (Latency)	首Token时间 + 生成时间	<500ms (首Token) <50ms/token (生成)
吞吐量 (Throughput)	每秒处理的tokens数	> 1000 tokens/s
并发数 (Concurrency)	同时处理的请求数	> 100
GPU利用率	GPU计算资源使用率	>80%

10.7.2 优化策略

```
class OptimizedLLMService:
    """
    优化的LLM服务
    """

    def __init__(self, model_path):
        # 1. 使用vLLM (最优推理引擎)
        self.llm = LLM(
            model=model_path,
            tensor_parallel_size=2,
            dtype="float16",
            gpu_memory_utilization=0.95, # 最大化GPU利用率
            max_num_seqs=256, # 最大并发序列数
            max_num_batched_tokens=8192, # 批处理token上限
        )

        # 2. 预热模型
        self._warmup()

        # 3. 启用缓存
        self.cache = {}

    def _warmup(self):
        """
        预热模型 (避免冷启动)
        """
        dummy_prompts = ["Hello"] * 10
        self.llm.generate(dummy_prompts, SamplingParams(max_tokens=10))
```

```

async def generate(self, prompt, params):
    """
    生成（带优化）
    """
    # 检查缓存
    cache_key = self._get_cache_key(prompt, params)
    if cache_key in self.cache:
        return self.cache[cache_key]

    # 生成
    result = await self._generate_with_retry(prompt, params)

    # 更新缓存
    self.cache[cache_key] = result

    return result

async def _generate_with_retry(self, prompt, params, max_retries=3):
    """
    带重试的生成
    """
    for attempt in range(max_retries):
        try:
            return self.llm.generate([prompt], SamplingParams(**params))
        except Exception as e:
            if attempt == max_retries - 1:
                raise
            await asyncio.sleep(0.1 * (2 ** attempt)) # 指数退避

```

10.8 成本优化

10.8.1 成本分析

```

class CostAnalyzer:
    """
    成本分析器
    """
    def __init__(self):
        self.gpu_cost_per_hour = {
            "A100-80GB": 3.0, # $/hour
            "A100-40GB": 2.5,
            "A10": 1.0,
            "T4": 0.5
        }

```

```

def calculate_cost(self, model_size, requests_per_day, avg_tokens_per_request):
    """
    计算日成本

    Args:
        model_size: 模型大小（如"7B", "13B", "70B"）
        requests_per_day: 每日请求数
        avg_tokens_per_request: 平均每请求生成tokens
    """
    # 1. 估算所需GPU
    gpu_type, num_gpus = self._estimate_gpu_requirement(model_size)

    # 2. 估算每请求时间
    time_per_request = avg_tokens_per_request * 0.05 # 假设50ms/token

    # 3. 计算每日GPU时间
    total_time_hours = (requests_per_day * time_per_request) / 3600

    # 4. 考虑利用率（通常只能达到50-70%）
    utilization = 0.6
    required_gpu_hours = (total_time_hours / utilization) * num_gpus

    # 5. 计算成本
    daily_cost = required_gpu_hours * self.gpu_cost_per_hour[gpu_type]

    return {
        "gpu_type": gpu_type,
        "num_gpus": num_gpus,
        "daily_cost": daily_cost,
        "monthly_cost": daily_cost * 30,
        "cost_per_1k_requests": daily_cost / (requests_per_day / 1000)
    }

def _estimate_gpu_requirement(self, model_size):
    """
    估算GPU需求
    """
    requirements = {
        "7B": ("A10", 1),
        "13B": ("A100-40GB", 1),
        "70B": ("A100-80GB", 4)
    }
    return requirements.get(model_size, ("A100-80GB", 1))

```

```
# 使用
analyzer = CostAnalyzer()
cost = analyzer.calculate_cost(
    model_size="7B",
    requests_per_day=100000,
    avg_tokens_per_request=200
)
print(f"Monthly cost: ${cost['monthly_cost']:.2f}")
```

10.8.2 成本优化策略

策略	节省	实现难度
量化 (INT8/INT4)	50-75%	☆☆
使用更小模型	70-90%	☆
Spot实例	60-70%	☆☆
批处理	30-50%	☆☆☆
缓存	20-40%	☆☆
请求去重	10-20%	☆

10.9 面试高频问题

Q1: 如何选择部署方案?

决策树:

请求量?

- └ <1000/day → API服务 (OpenAI等)
- └ 1000-10000/day → 云GPU实例 + vLLM
- └ >10000/day → 自建集群 + K8s

数据敏感性?

- └ 高 → 自建/私有云
- └ 低 → 公有云

延迟要求?

- └ <100ms → 边缘部署
- └ <1s → 标准云部署

Q2: vLLM为什么快?

核心技术:

1. **PagedAttention**: KV Cache分页管理
2. **Continuous Batching**: 动态组batch
3. **优化的CUDA kernels**: 高效的底层实现

性能对比:

HuggingFace: 100 tokens/s

vLLM: 800+ tokens/s (8x提升)

Q3: 如何处理突发流量?

策略:

1. **自动扩容**: K8s HPA
2. **限流**: 保护后端
3. **队列**: 削峰填谷
4. **降级**: 返回cached或简化回复

限流+队列组合

```
@app.post("/v1/completions")
```

```
@limiter.limit("100/minute")
```

```
async def create_completion(request: CompletionRequest):
```

```
    if active_requests.get() > MAX_CONCURRENT:
```

```
        # 超过并发上限, 加入队列
```

```
        return await enqueue_request(request)
```

```
    else:
```

```
        return await process_immediately(request)
```

Q4: 如何降低推理成本?

优先级排序:

1. **量化** (最高ROI) : 成本降50%+, 精度损失<5%
2. **批处理**: 吞吐量提升3-5x
3. **缓存**: 重复请求免费
4. **更小模型**: 7B vs 70B, 成本差10倍

Q5: 生产环境必备的监控指标?

三大类:

1. 业务指标:

- QPS (每秒请求数)
- 延迟 (P50, P95, P99)
- 错误率

2. 系统指标:

- GPU利用率
- GPU显存
- CPU/内存

3. 成本指标:

- 每1K请求成本
- GPU空闲时间
- 缓存命中率

10.10 本章小结

本章全面介绍了大模型的部署与服务化:

✅ **部署方案**: 本地、云服务、自建集群 ✅ **推理引擎**: vLLM是首选 ✅ **API设计**: FastAPI + 异步 + 批处理 ✅ **K8s部署**: 容器化 + 自动扩缩容 ✅ **监控日志**: Prometheus + 结构化日志 ✅ **成本优化**: 量化、批处理、缓存

关键点:

- vLLM是生产环境首选推理引擎
- 批处理和缓存能显著提升吞吐量
- 监控和日志对生产至关重要
- 量化是成本优化的最佳手段

下一章预告: 第11章将讲解模型评估方法和基准测试。