

第12章 模型优化与调试

■ 优化是一门艺术，调试是一门科学

12.1 训练问题诊断

12.1.1 Loss曲线分析

```
import matplotlib.pyplot as plt

def analyze_training_loss(loss_history):
    """
    分析训练loss曲线
    """
    plt.figure(figsize=(12, 4))

    # 绘制Loss曲线
    plt.subplot(1, 2, 1)
    plt.plot(loss_history['train_loss'], label='Train Loss')
    plt.plot(loss_history['val_loss'], label='Val Loss')
    plt.xlabel('Step')
    plt.ylabel('Loss')
    plt.legend()
    plt.title('Loss Curve')

    # 平滑后的曲线
    plt.subplot(1, 2, 2)
    smooth_train = moving_average(loss_history['train_loss'], window=100)
    smooth_val = moving_average(loss_history['val_loss'], window=100)
    plt.plot(smooth_train, label='Train Loss (smoothed)')
    plt.plot(smooth_val, label='Val Loss (smoothed)')
    plt.xlabel('Step')
    plt.ylabel('Loss')
    plt.legend()
    plt.title('Smoothed Loss Curve')

    plt.tight_layout()
    plt.savefig('loss_analysis.png')

    # 诊断
    diagnosis = diagnose_loss_curve(loss_history)
    return diagnosis

def diagnose_loss_curve(loss_history):
    """
```

诊断loss曲线问题

"""

```
train_loss = loss_history['train_loss']  
val_loss = loss_history['val_loss']
```

```
issues = []
```

1. 检查是否不下降

```
if train_loss[-1] > train_loss[0] * 0.9:  
    issues.append({  
        "problem": "训练loss不下降",  
        "possible_causes": [  
            "学习率过小",  
            "学习率过大（爆炸）",  
            "数据问题",  
            "模型初始化问题"  
        ],  
        "solutions": [  
            "调整学习率（尝试1e-5到1e-3）",  
            "检查梯度范数",  
            "验证数据质量",  
            "使用更好的初始化"  
        ]  
    })
```

2. 检查过拟合

```
gap = val_loss[-1] - train_loss[-1]  
if gap > train_loss[-1] * 0.3:  
    issues.append({  
        "problem": "过拟合",  
        "possible_causes": [  
            "模型太大",  
            "训练数据太少",  
            "训练时间太长"  
        ],  
        "solutions": [  
            "增加正则化（dropout, weight decay）",  
            "数据增强",  
            "Early stopping",  
            "使用更小的模型"  
        ]  
    })
```

3. 检查欠拟合

```
if train_loss[-1] > expected_final_loss * 1.5:  
    issues.append({  
        "problem": "欠拟合",  
        "possible_causes": [  
            "模型太小",  
            "训练数据太少",  
            "训练时间太短"  
        ],  
        "solutions": [  
            "增加模型复杂度",  
            "增加训练数据",  
            "增加训练时间"  
        ]  
    })
```

```

        "训练不充分",
        "学习率太小"
    ],
    "solutions": [
        "增大模型",
        "训练更多epochs",
        "提高学习率",
        "检查数据质量"
    ]
})

# 4. 检查震荡
volatility = np.std(train_loss[-100:]) / np.mean(train_loss[-100:])
if volatility > 0.1:
    issues.append({
        "problem": "训练不稳定/震荡",
        "possible_causes": [
            "学习率过大",
            "batch size太小",
            "梯度爆炸"
        ],
        "solutions": [
            "降低学习率",
            "增大batch size",
            "梯度裁剪",
            "使用warmup"
        ]
    })

if not issues:
    return {"status": "健康", "message": "训练曲线看起来正常"}

return {"status": "有问题", "issues": issues}

```

12.1.2 梯度监控

```

class GradientMonitor:
    """
    梯度监控器
    """
    def __init__(self, model):
        self.model = model
        self.gradient_norms = []
        self.param_norms = []

    def log_gradients(self):
        """
        记录梯度信息

```

```

"""
total_norm = 0
param_norm = 0

for name, param in self.model.named_parameters():
    if param.grad is not None:
        # 参数梯度的范数
        grad_norm = param.grad.data.norm(2).item()
        total_norm += grad_norm ** 2

        # 参数本身的范数
        param_norm += param.data.norm(2).item() ** 2

total_norm = total_norm ** 0.5
param_norm = param_norm ** 0.5

self.gradient_norms.append(total_norm)
self.param_norms.append(param_norm)

return total_norm, param_norm

def check_gradient_health(self):
    """
    检查梯度健康状况
    """
    if not self.gradient_norms:
        return "No gradient data"

    recent_grads = self.gradient_norms[-100:]
    avg_grad = np.mean(recent_grads)

    issues = []

    # 1. 梯度消失
    if avg_grad < 1e-7:
        issues.append({
            "problem": "梯度消失",
            "solutions": [
                "检查激活函数（避免sigmoid）",
                "使用残差连接",
                "使用Layer Normalization",
                "降低网络深度"
            ]
        })

    # 2. 梯度爆炸
    if avg_grad > 100:
        issues.append({
            "problem": "梯度爆炸",

```

```

        "solutions": [
            "梯度裁剪 (clip_grad_norm) ",
            "降低学习率",
            "检查数据是否异常",
            "使用Layer Normalization"
        ]
    })

```

3. 梯度NaN

```

if np.isnan(recent_grads).any():
    issues.append({
        "problem": "梯度包含NaN",
        "solutions": [
            "检查数据是否有NaN/Inf",
            "降低学习率",
            "使用混合精度时注意loss scaling",
            "检查除零错误"
        ]
    })

```

```

return issues if issues else "健康"

```

```

def plot_gradient_flow(self):
    """
    可视化梯度流
    """
    plt.figure(figsize=(12, 5))

    # 绘制各层梯度
    layer_grads = {}
    for name, param in self.model.named_parameters():
        if param.grad is not None:
            layer_name = name.split('.')[0]
            grad_norm = param.grad.data.norm(2).item()

            if layer_name not in layer_grads:
                layer_grads[layer_name] = []
            layer_grads[layer_name].append(grad_norm)

    # 绘制
    layers = list(layer_grads.keys())
    avg_grads = [np.mean(layer_grads[l]) for l in layers]

    plt.bar(range(len(layers)), avg_grads)
    plt.xticks(range(len(layers)), layers, rotation=45)
    plt.ylabel('Average Gradient Norm')
    plt.title('Gradient Flow Across Layers')

```

```
plt.tight_layout()
plt.savefig('gradient_flow.png')
```

12.1.3 学习率调优

```
def learning_rate_finder(model, train_loader, optimizer, criterion):
    """
    学习率范围测试 (LR Finder)

    在不同学习率下训练，找到最佳学习率范围
    """
    model.train()

    # 学习率范围: 1e-7 到 1
    lr_min = 1e-7
    lr_max = 1
    num_steps = 100

    lrs = np.logspace(np.log10(lr_min), np.log10(lr_max), num_steps)
    losses = []

    for i, lr in enumerate(lrs):
        # 设置学习率
        for param_group in optimizer.param_groups:
            param_group['lr'] = lr

        # 训练一步
        batch = next(iter(train_loader))
        optimizer.zero_grad()

        outputs = model(batch['input_ids'])
        loss = criterion(outputs, batch['labels'])

        loss.backward()
        optimizer.step()

        losses.append(loss.item())

        # 如果Loss爆炸，提前停止
        if loss.item() > losses[0] * 10:
            break

    # 绘制结果
    plt.figure(figsize=(10, 5))
    plt.semilogx(lrs[:len(losses)], losses)
    plt.xlabel('Learning Rate')
    plt.ylabel('Loss')
    plt.title('Learning Rate Finder')
```

```

plt.grid(True)
plt.savefig('lr_finder.png')

# 找到最佳学习率 (Loss下降最快的点)
gradients = np.gradient(losses)
best_lr_idx = np.argmin(gradients)
best_lr = lrs[best_lr_idx]

print(f"建议学习率: {best_lr:.2e}")
print(f"建议范围: {best_lr/10:.2e} 到 {best_lr*2:.2e}")

return best_lr

```

12.2 数据问题排查

12.2.1 数据质量检查

```

class DataQualityChecker:
    """
    数据质量检查器
    """

    def __init__(self, dataset, tokenizer):
        self.dataset = dataset
        self.tokenizer = tokenizer

    def run_checks(self):
        """
        运行所有检查
        """
        print("检查数据质量...")

        results = {}

        results['length_stats'] = self.check_length_distribution()
        results['duplicates'] = self.check_duplicates()
        results['quality'] = self.check_text_quality()
        results['balance'] = self.check_class_balance()

        self.generate_report(results)

        return results

    def check_length_distribution(self):
        """
        检查长度分布
        """
        lengths = []

```

```

for example in self.dataset:
    text = example['text']
    tokens = self.tokenizer.encode(text)
    lengths.append(len(tokens))

stats = {
    "mean": np.mean(lengths),
    "std": np.std(lengths),
    "min": np.min(lengths),
    "max": np.max(lengths),
    "median": np.median(lengths),
    "p95": np.percentile(lengths, 95)
}

# 绘制分布
plt.figure(figsize=(10, 5))
plt.hist(lengths, bins=50)
plt.xlabel('Length (tokens)')
plt.ylabel('Count')
plt.title('Length Distribution')
plt.axvline(stats['mean'], color='r', linestyle='--', label=f"Mean: {stats['mean']}")
plt.legend()
plt.savefig('length_distribution.png')

# 检查异常
warnings = []
if stats['std'] > stats['mean']:
    warnings.append("长度方差很大，考虑分桶或裁剪")

if stats['max'] > 2048:
    warnings.append(f"存在超长样本（{stats['max']} tokens），可能需要截断")

return {"stats": stats, "warnings": warnings}

def check_duplicates(self):
    """
    检查重复数据
    """
    from collections import Counter

    texts = [example['text'] for example in self.dataset]

    # 完全重复
    exact_duplicates = len(texts) - len(set(texts))

    # 近似重复（使用MinHash）
    # ... 实现略

```



```

duplicate_rate = exact_duplicates / len(texts)

if duplicate_rate > 0.01:
    warning = f"检测到{duplicate_rate:.1%}的重复数据，建议去重"
else:
    warning = None

return {
    "exact_duplicates": exact_duplicates,
    "duplicate_rate": duplicate_rate,
    "warning": warning
}

def check_text_quality(self):
    """
    检查文本质量
    """
    issues = []

    for i, example in enumerate(self.dataset[:1000]): # 抽样检查
        text = example['text']

        # 检查1: 过短
        if len(text) < 10:
            issues.append(f"样本{i}: 文本过短")

        # 检查2: 特殊字符比例过高
        special_char_ratio = sum(not c.isalnum() and not c.isspace()
                                for c in text) / len(text)
        if special_char_ratio > 0.3:
            issues.append(f"样本{i}: 特殊字符过多")

        # 检查3: 重复字符
        if any(text.count(c*10) > 0 for c in 'abcdefghijklmnopqrstuvwxyz'):
            issues.append(f"样本{i}: 检测到重复字符")

    return {
        "issues_found": len(issues),
        "sample_issues": issues[:10] # 只显示前10个
    }

def check_class_balance(self):
    """
    检查类别平衡（如果是分类任务）
    """
    if 'label' not in self.dataset[0]:
        return "N/A (not a classification task)"

from collections import Counter

```

```

labels = [example['label'] for example in self.dataset]
label_counts = Counter(labels)

# 计算不平衡程度
max_count = max(label_counts.values())
min_count = min(label_counts.values())
imbalance_ratio = max_count / min_count

warning = None
if imbalance_ratio > 10:
    warning = f"类别严重不平衡 ({imbalance_ratio:.1f}:1)，考虑重采样或加权"

return {
    "label_distribution": dict(label_counts),
    "imbalance_ratio": imbalance_ratio,
    "warning": warning
}

```

12.2.2 数据增强

```

def augment_data(text):
    """
    数据增强
    """
    augmented = []

    # 1. 回译 (Back-translation)
    # translated = translate(text, 'en', 'fr')
    # back_translated = translate(translated, 'fr', 'en')
    # augmented.append(back_translated)

    # 2. 同义词替换
    # augmented.append(synonym_replacement(text))

    # 3. 随机插入
    # augmented.append(random_insertion(text))

    # 4. 随机删除
    # augmented.append(random_deletion(text, p=0.1))

    # 5. 随机交换
    # augmented.append(random_swap(text, n=3))

    return augmented

```

12.3 性能优化

12.3.1 训练加速技巧

技巧1: 梯度累积

```
def train_with_gradient_accumulation(model, dataloader, accumulation_steps=4):  
    """  
    梯度累积: 模拟大batch size  
    """  
  
    optimizer.zero_grad()  
  
    for i, batch in enumerate(dataloader):  
        # 前向+反向  
        loss = model(batch)  
        loss = loss / accumulation_steps # 归一化  
        loss.backward()  
  
        # 每accumulation_steps步更新一次  
        if (i + 1) % accumulation_steps == 0:  
            optimizer.step()  
            optimizer.zero_grad()
```

技巧2: 混合精度训练

```
from torch.cuda.amp import autocast, GradScaler  
  
scaler = GradScaler()  
  
for batch in dataloader:  
    optimizer.zero_grad()  
  
    # 使用autocast  
    with autocast():  
        output = model(batch)  
        loss = criterion(output, labels)  
  
    # scaled backward  
    scaler.scale(loss).backward()  
    scaler.step(optimizer)  
    scaler.update()
```

技巧3: DataLoader优化

```
train_loader = DataLoader(  
    dataset,  
    batch_size=32,  
    num_workers=4, # 多进程加载  
    pin_memory=True, # 锁页内存  
    prefetch_factor=2, # 预取  
)
```

技巧4: 编译模型 (PyTorch 2.0+)

```
model = torch.compile(model)
```

12.3.2 显存优化

技巧1: 梯度检查点 (Gradient Checkpointing)

```
from torch.utils.checkpoint import checkpoint
```

```
class TransformerBlockWithCheckpoint(nn.Module):
```

```
    def forward(self, x):
```

```
        # 使用checkpoint减少显存
```

```
        return checkpoint(self._forward, x)
```

```
    def _forward(self, x):
```

```
        x = self.attention(x)
```

```
        x = self.ffn(x)
```

```
        return x
```

技巧2: 清理缓存

```
import gc
```

```
import torch
```

```
def clear_memory():
```

```
    """
```

```
    清理显存
```

```
    """
```

```
    gc.collect()
```

```
    torch.cuda.empty_cache()
```

技巧3: 使用更小的数据类型

```
model = model.half() # FP16
```

或

```
model = model.bfloat16() # BF16
```

技巧4: 零冗余优化器 (ZeRO)

```
from deepspeed import DeepSpeedEngine
```

```
model_engine = DeepSpeedEngine(
```

```
    model=model,
```

```
    config={
```

```
        "zero_optimization": {
```

```
            "stage": 3, # ZeRO Stage 3
```

```
        }
```

```
}  
)
```

12.4 超参数调优

12.4.1 网格搜索 vs 随机搜索

```
from sklearn.model_selection import ParameterGrid  
import random  
  
# 网格搜索  
param_grid = {  
    'learning_rate': [1e-5, 5e-5, 1e-4],  
    'batch_size': [16, 32, 64],  
    'warmup_steps': [100, 500, 1000]  
}  
  
grid = ParameterGrid(param_grid)  
  
best_score = 0  
best_params = None  
  
for params in grid:  
    score = train_and_evaluate(model, params)  
    if score > best_score:  
        best_score = score  
        best_params = params  
  
print(f"Best params: {best_params}")  
print(f"Best score: {best_score}")  
  
# 随机搜索（更高效）  
def random_search(param_distributions, n_iter=10):  
    """  
    随机搜索超参数  
    """  
    results = []  
  
    for i in range(n_iter):  
        # 随机采样超参数  
        params = {  
            'learning_rate': random.choice([1e-5, 5e-5, 1e-4, 5e-4]),  
            'batch_size': random.choice([16, 32, 64]),  
            'warmup_steps': random.randint(100, 1000),  
            'weight_decay': random.uniform(0.01, 0.1)  
        }
```

```

# 训练评估
score = train_and_evaluate(model, params)

results.append({
    'params': params,
    'score': score
})

# 排序
results.sort(key=lambda x: x['score'], reverse=True)

return results[0] # 返回最佳结果

```

12.4.2 贝叶斯优化

```

from skopt import gp_minimize
from skopt.space import Real, Integer, Categorical

# 定义搜索空间
space = [
    Real(1e-6, 1e-3, name='learning_rate', prior='log-uniform'),
    Integer(8, 64, name='batch_size'),
    Integer(100, 2000, name='warmup_steps'),
    Real(0.0, 0.2, name='dropout'),
]

def objective(params):
    """
    目标函数（最小化）
    """
    lr, batch_size, warmup_steps, dropout = params

    # 训练模型
    val_loss = train_model(
        learning_rate=lr,
        batch_size=batch_size,
        warmup_steps=warmup_steps,
        dropout=dropout
    )

    return val_loss # 返回验证损失

# 运行贝叶斯优化
result = gp_minimize(
    objective,
    space,
    n_calls=20, # 评估20组参数

```

```

        random_state=42
    )

    print(f"Best parameters: {result.x}")
    print(f"Best validation loss: {result.fun}")

```

12.5 调试技巧

12.5.1 单样本过拟合测试

```

def sanity_check_overfit_single_batch(model, single_batch, num_steps=100):
    """
    Sanity check: 在单个batch上过拟合

    如果模型无法在单个batch上过拟合，说明有bug
    """
    model.train()
    optimizer = torch.optim.Adam(model.parameters(), lr=1e-3)

    losses = []

    for step in range(num_steps):
        optimizer.zero_grad()

        outputs = model(single_batch['input_ids'])
        loss = criterion(outputs, single_batch['labels'])

        loss.backward()
        optimizer.step()

        losses.append(loss.item())

        if step % 10 == 0:
            print(f"Step {step}, Loss: {loss.item():.4f}")

    # 检查是否成功过拟合
    if losses[-1] < losses[0] * 0.1:
        print("✓ 通过: 模型能够在单个batch上过拟合")
    else:
        print("X 失败: 模型无法过拟合，可能存在问题")
        print("  - 检查损失函数")
        print("  - 检查数据标签")
        print("  - 检查模型前向传播")

    return losses

```

12.5.2 激活值和权重检查

```

def check_activations_and_weights(model, batch):
    """
    检查激活值和权重
    """

    model.eval()

    # 注册hook收集激活值
    activations = {}

    def get_activation(name):
        def hook(module, input, output):
            activations[name] = output.detach()
        return hook

    # 为每层注册hook
    for name, module in model.named_modules():
        module.register_forward_hook(get_activation(name))

    # 前向传播
    with torch.no_grad():
        _ = model(batch['input_ids'])

    # 检查
    issues = []

    for name, activation in activations.items():
        # 检查NaN
        if torch.isnan(activation).any():
            issues.append(f"{name}: 包含NaN")

        # 检查Inf
        if torch.isinf(activation).any():
            issues.append(f"{name}: 包含Inf")

        # 检查全零
        if (activation == 0).all():
            issues.append(f"{name}: 全零激活 (可能是dead neurons)")

        # 检查范围
        act_max = activation.abs().max().item()
        if act_max > 1000:
            issues.append(f"{name}: 激活值过大 ({act_max})")

    # 检查权重
    for name, param in model.named_parameters():
        # 检查NaN
        if torch.isnan(param).any():
            issues.append(f"{name}: 权重包含NaN")

```



```

# 检查权重范围
weight_std = param.std().item()
if weight_std < 1e-5:
    issues.append(f"{name}: 权重方差过小 ({weight_std}) ")
elif weight_std > 10:
    issues.append(f"{name}: 权重方差过大 ({weight_std}) ")

if issues:
    print("发现问题:")
    for issue in issues:
        print(f" - {issue}")
else:
    print("✓ 激活值和权重检查通过")

return activations, issues

```

12.5.3 常见错误排查清单

```

debugging_checklist = {
    "Loss不下降": [
        "□ 检查学习率（太大或太小）",
        "□ 检查数据是否shuffle",
        "□ 检查标签是否正确",
        "□ 检查损失函数是否合适",
        "□ 尝试更简单的模型",
        "□ 检查输入是否normalize",
    ],
    "Loss变NaN": [
        "□ 降低学习率",
        "□ 添加梯度裁剪",
        "□ 检查数据中是否有NaN/Inf",
        "□ 使用Layer Normalization",
        "□ 检查除零错误",
        "□ 使用更稳定的数据类型（FP32）",
    ],
    "显存不足": [
        "□ 减小batch size",
        "□ 使用梯度累积",
        "□ 启用梯度检查点",
        "□ 使用混合精度",
        "□ 减小模型大小",
        "□ 使用ZeRO优化器",
    ],
    "训练太慢": [

```

```

        "□ 增大batch size",
        "□ 使用多GPU训练",
        "□ 优化DataLoader (num_workers) ",
        "□ 使用混合精度",
        "□ 编译模型 (torch.compile) ",
        "□ 检查是否有瓶颈操作",
    ],

    "过拟合": [
        "□ 增加Dropout",
        "□ 增加权重衰减",
        "□ 数据增强",
        "□ Early stopping",
        "□ 减小模型",
        "□ 增加训练数据",
    ]
}

def print_debugging_checklist(problem):
    """
    打印调试清单
    """
    if problem in debugging_checklist:
        print(f"\n{problem}排查清单: ")
        for item in debugging_checklist[problem]:
            print(f"    {item}")
    else:
        print("未找到对应问题的清单")

```

12.6 本章小结

本章介绍了模型优化与调试的实用技巧：

✅ **训练诊断**：Loss曲线分析、梯度监控、学习率调优
 ✅ **数据排查**：质量检查、去重、增强
 ✅ **性能优化**：训练加速、显存优化
 ✅ **超参调优**：网格搜索、贝叶斯优化
 ✅ **调试技巧**：单样本过拟合、激活值检查

关键点：

- 先检查数据，再怀疑模型
- 从简单开始调试（单样本过拟合）
- 监控梯度和激活值
- 系统化超参数搜索

完成！ 至此第四部分（评估与优化篇）全部完成。

下一部分预告： 第五部分将介绍领域应用，第六部分补充系统设计和公司分析。